

Work1:

Model:1 GPT – Provide -OpenAI – Useful for Text to Speeac/speech to text

OpenAI (GPT / o-series)

Model	Knowledge Cutoff	Internet Access	Notes
GPT-5.2	August 2025	Yes (browsing)	Latest flagship
GPT-5.1	September 2024	Yes (browsing)	Previous generation
GPT-5	October 2024	Yes (browsing)	First GPT-5 release
GPT-5 mini	May 2024	Yes (browsing)	Lightweight variant
GPT-4.1	June 2024	Yes (browsing)	Mid-cycle update
GPT-4o	October 2023	Yes (browsing)	Still widely used in integrations
GPT-4o mini	October 2023	Yes (browsing)	Cost-efficient variant
o4-mini	June 2024	Yes	Latest reasoning model
o3 / o3-pro	June 2024	Yes	Advanced reasoning
o1 / o1-pro	October 2023	Yes	First-gen reasoning models

Model : 2 Gemini – provider – Google – Text to image/ Video generation


The screenshot shows the OpenRouter website with the URL `openrouter.ai/provider/google-vertex`. The page lists several Google Gemini models:

- Google: Gemini 3.1 Flash Lite**: Gemini 3.1 Flash Lite is Google's GA high-efficiency multimodal model optimized for low-latency, high-volume workloads. It supports text, image, video, audio, and PDF inputs, and is designed for lightweight ...
by google · May 7, 2026 · 1.05M context · \$0.25/M input tokens · \$1.50/M output tokens
- Google: Chirp 3**: Chirp 3 is Google's latest multilingual speech-to-text model. It offers enhanced transcription accuracy across 24 GA languages and 77+ preview languages, with support for automatic language detection, ...
by google · May 5, 2026 · \$0.016 per minute
- Google: Gemini 3.1 Flash TTS Preview**: Gemini 3.1 Flash TTS Preview is a text-to-speech model from Google, and a substantial generational step up from Gemini 2.5 Flash TTS. It takes text input and produces audio output across 70+ languages — nearl...
by google · Apr 24, 2026 · 8K context · \$1/M input tokens · \$20/M output tokens
- Google: Veo 3.1 Fast**: Google's mid-tier video generation model balancing speed and quality. Veo 3.1 Fast generates high-quality video from text or image prompts with native synchronized audio, offering faster turnaround than Veo 3.1 a...
by google · Apr 24, 2026 · from \$0.10 per second
- Google: Veo 3.1 Lite**: Google's most cost-effective video generation model, designed for high-volume applications and rapid iteration. Veo 3.1 Lite generates 720p and 1080p video from text or image prompts with native synchroniz...

The screenshot shows the OpenRouter website with the URL `openrouter.ai/provider/google-vertex`. The page lists several Google Gemini models:

- Google: Gemma 4 26B A4B**: Gemma 4 26B A4B IT is an instruction-tuned Mixture-of-Experts (MoE) model from Google DeepMind. Despite 25.2B total parameters, only 3.8B activate per token during inference — delivering near-31B qualit...
by google · Apr 3, 2026 · 262K context · \$0.15/M input tokens · \$0.60/M output tokens
- Google: Veo 3.1**: Google's state-of-the-art video generation model, built for maximum visual fidelity in final production cuts. Veo 3.1 generates high-quality 1080p video from text or image prompts with native synchronized audio —...
by google · Mar 23, 2026 · from \$0.40 per second
- Google: Gemini 3.1 Flash Lite Preview**: Gemini 3.1 Flash Lite Preview is Google's high-efficiency model optimized for high-volume use cases. It outperforms Gemini 2.5 Flash Lite on overall quality and approaches Gemini 2.5 Flash performance across ...
by google · Mar 3, 2026 · 1.05M context · \$0.25/M input tokens · \$1.50/M output tokens
- Google: Nano Banana 2 (Gemini 3.1 Flash Image Preview)**: Gemini 3.1 Flash Image Preview, a.k.a. "Nano Banana 2," is Google's latest state of the art image generation and editing model, delivering Pro-level visual quality at Flash speed. It combines advanced contextual ...
by google · Feb 26, 2026 · 131K context · \$0.50/M input tokens · \$3/M output tokens
- Google: Gemini 3.1 Pro Preview**: Gemini 3.1 Pro Preview is Google's frontier reasoning model, delivering enhanced software engineering performance, improved agentic reliability, and more efficient token usage across complex workflows. ...
by google · Feb 19, 2026 · 1.05M context · \$2/M input tokens · \$12/M output tokens

Model 3: Claude -Provider-Antropic - Useful for Code generation



Model: 4 Qwen -provider-Alibaba -coverts text, image, video to Text and reasoning task

links

Conti

post

Gro

Gro

gma

Inbo

Recr

Wha

(2) V

ram

cut c

AI Ki

Chai

Ope

Ask Gemini

openrouter.ai/qwen/qwen3.6-35b-a3b

Campus Portal:Tata... jobs WORK hctc quick Drop me in Bookma... ruby automation TRADE Automation Java simple Recipes Gmail YouTube Maps All Bookmarks

OpenRouter

Search

May 2026

Recraft: Recraft V4 Pro

Recraft: Recraft V4

Recraft: Recraft V3

Google: Gemini 3.1 Flash Lite

Baidu Qianfan: CoBuddy (free)

OpenAI: GPT Chat Latest

Google: Chirp 3

OpenAI: GPT-4o Mini Transcribe

wen3.6 35B A3B

Chat

Compare

6-35b-a3b

2026 262,144 context \$0.15/M input tokens \$1/M output tokens

A3B is an open-weight multimodal model from Alibaba Cloud with 35 billion total parameters and 3 billion active parameters per token. It uses a hybrid sparse mixture-of-experts architecture combining Gated DeltaNet linear attention with standard gated attention layers, enabling efficient inference at a fraction of the compute cost. The model supports a 262K token native context window (extensible to 1M via YaRN) and accepts text, image, and video inputs. It includes integrated thinking mode with reasoning traces preserved across multi-turn conversations, function calling, and structured output. Released under the Apache 2.0 license.

Standard

Model weights

Overview

Playground

Providers

Performance

Pricing

Apps

Activity

Uptime

API

Providers for Qwen3.6 35B A3B

OpenRouter routes requests to the best providers that are able to handle your prompt size and parameters, with fallbacks to maximize uptime.

Sort by

Parasail

Latency

Throughput

Uptime

0.89s

64tps

Qwen: Qwen3.5 Plus 2026-04-20

Chat

Compare

qwen/qwen3.5-plus-20260420

Released Apr 27, 2026 | 1,000,000 context | \$0.40/M input tokens | \$2.40/M output tokens

Qwen3.5 Plus (April 2026) is a large-scale multimodal language model from Alibaba. It accepts text, image, and video input and produces text output, with a 1M token context window. This is an updated version of Qwen3.5 Plus with tiered pricing above 256K tokens.

Overview

Playground

Providers

Performance

Pricing

Apps

Activity

Uptime

API

Providers for Qwen3.5 Plus 2026-04-20

OpenRouter [routes requests](#) to the best providers that are able to handle your prompt size and parameters, with fallbacks to maximize uptime.

Alibaba Cloud Int.				Latency	Throughput	Uptime
50				0.70s	11tps	
Total Context	Max Output	Input Price	Output Price			
1M	65.5K	<256K \$0.40 >256K \$0.50 /K Tokens	<256K \$2.40 >256K \$3 /K Tokens			

Qwen: Qwen3.6 Max Preview

Chat

Compare

qwen/qwen3.6-max-preview

Released Apr 27, 2026 | 262,144 context | \$1.04/M input tokens | \$6.24/M output tokens

[Finance \(#49\)](#)

Qwen3.6-Max-Preview is a proprietary frontier model from Alibaba Cloud built on a sparse mixture-of-experts architecture with approximately 1 trillion total parameters. It is optimized for agentic coding, tool use, and long-context reasoning, supporting a 262K token context window. The model includes an integrated thinking mode that preserves reasoning traces across multi-turn conversations and supports structured output and function calling. Access is available exclusively through the Alibaba Cloud Model Studio and Qwen Studio APIs; no open weights are provided.

Overview

Playground

Providers

Performance

Pricing

Apps

Activity

Uptime

API

Providers for Qwen3.6 Max Preview

OpenRouter [routes requests](#) to the best providers that are able to handle your prompt size and parameters, with fallbacks to maximize uptime.

Alibaba Cloud Int.				Latency	Throughput	Uptime
50				1.45s	34tps	
Total Context	Max Output	Input Price	Output Price	Cache Creation (Server)		
262.1K	65.5K	<128K \$1-30 \$1.04 >128K \$2 \$1.60	<128K \$7-60 \$6.24 >128K \$2 \$9.60	<128K \$1-60 \$1.30 >128K \$2 \$2		

openrouter.ai/qwen/qwen3.6-27b

Qwen: Qwen3.6 27B

Released Apr 27, 2026 262,144 context \$0.32/M input tokens \$3.20/M output tokens

SEO (#22) Programming (#48)

Qwen3.6 27B is a dense 27-billion-parameter language model from the Qwen Team at Alibaba, released in April 2026. It features hybrid multimodal capabilities — accepting text, image, and video inputs — and supports a 262,144-token context window. The model is designed for agentic coding and reasoning tasks, with particular strength in repository-level code comprehension, front-end development workflows, and multi-step problem solving. It includes a built-in thinking mode for extended reasoning and preserves thinking context across conversation history. Qwen3.6 27B supports 201 languages and dialects and is released under the Apache 2.0 license.

Standard Model weights

Overview Playground Providers Performance Pricing Apps Activity Uptime API

Providers for Qwen3.6 27B

OpenRouter routes requests to the best providers that are able to handle your prompt size and parameters, with fallbacks to maximize uptime.

Filter quantization Sort by

Provider	Latency	Throughput	Uptime
DeepInfra	0.35s	44tps	100%

openrouter.ai/qwen/qwen3.5-35b-a3b

Qwen: Qwen3.5-35B-A3B

Released Feb 25, 2026 262,144 context \$0.14/M input tokens \$1/M output tokens

Health (#31) Programming (#41)

The Qwen3.5 Series 35B-A3B is a native vision-language model designed with a hybrid architecture that integrates linear attention mechanisms and a sparse mixture-of-experts model, achieving higher inference efficiency. Its overall performance is comparable to that of the Qwen3.5-72B.

Standard Model weights

Overview Playground Providers Performance Pricing Apps Activity Uptime API

Providers for Qwen3.5-35B-A3B

OpenRouter routes requests to the best providers that are able to handle your prompt size and parameters, with fallbacks to maximize uptime.

Filter quantization Sort by

Provider	Latency	Throughput	Uptime
DeepInfra	0.33s	57tps	100%

Total Context	Max Output	Input Price	Output Price	Cache Read
16384	128	\$0.14	\$1.00	\$0.10

Model: 5 Deepseek-Provider-Deepseek – good in reasoning

link

Con

pos

Gro

Gro

gms

lnb

R11

Rec

whi

(2)

ram

cut

Al

Chi

dee


Ask Gemini

openrouter.ai/provider/deepseek

Campus Portal:Tata...jobsWORKhrctc quickDrop me in Bookma...ruby automationTRADEAutomationJava simpleRecipesGmailYouTubeMapsAll Bookmarks

OpenRouterSearch

FusionModelsChatRankingsAppsEnterprisePricingDocsSign Up



80B
40B

Feb 8Feb 18Feb 28Mar 10Mar 20Mar 30Apr 9Apr 19Apr 29

DeepSeek: DeepSeek V4 Pro

DeepSeek V4 Pro is a large-scale Mixture-of-Experts model from DeepSeek with 1.6T total parameters and 49B activated parameters, supporting a 1M-token context window. It is designed for advanced reasoning, ...
by [deepseek](#) · Apr 24, 2026 · 1.05M context · \$0.435/M input tokens · \$0.87/M output tokens

DeepSeek: DeepSeek V4 Flash

DeepSeek V4 Flash is an efficiency-optimized Mixture-of-Experts model from DeepSeek with 284B total parameters and 1.3B activated parameters, supporting a 1M-token context window. It is designed for fast ...
by [deepseek](#) · Apr 24, 2026 · 1.05M context · \$0.14/M input tokens · \$0.28/M output tokens

OpenRouter

© 2026 OpenRouter, Inc.

Product

Chat

Company

About

Developer

Documentation

Connect

Discord

link | Cor | pos | Gro | Gro | gms | Inb | R1 | Rec | whi | (2) | ram | cut | AI | Chi | dee | + | Ask Gemini

openrouter.ai/deepseek/deepseek-v3.2

Campus Portal:Tata... | jobs | WORK | iktc quick | Drop me in Bookma... | ruby automation | TRADE | Automation | Java simple | Recipes | Gmail | YouTube | Maps | All Bookmarks

OpenRouter | Search | / | Fusion | Models | Chat | Rankings | Apps | Enterprise | Pricing | Docs | Sign Up

Academia (#8) | Finance (#19) | Health (#6) | Legal (#15) | Marketing (#20) | +7 categories

DeepSeek-V3.2 is a large language model designed to harmonize high computational efficiency with strong reasoning and agentic tool-use performance. It introduces DeepSeek Sparse Attention (DSA), a fine-grained sparse attention mechanism that reduces training and inference cost.

Standard | Model weights

Overview | Playground | Providers | Performance | Pricing | Apps | Activity | Uptime | API

Providers for DeepSeek V3.2

OpenRouter routes requests to the best providers that are able to handle your prompt size and parameters, with fallbacks to maximize uptime.

Filter quantization | Sort by

Baidu Qianfan

Latency 6.33s | Throughput 21tps | Uptime

Total Context	Max Output	Input Price	Output Price	Cache Read
131.1K	65.5K	\$0.252 /K Tokens	\$0.378 /K Tokens	\$0.0252 /K Tokens

SiliconFlow

Latency 2.62s | Throughput 20tps | Uptime

Total Context	Max Output	Input Price	Output Price	Cache Read
163.8K	163.8K	\$0.259 /K Tokens	\$0.42 /K Tokens	\$0.135 /K Tokens

link | Cor | pos | Gro | Gro | gms | Inb | R1 | Rec | whi | (2) | ram | cut | AI | Chi | dee | + | Ask Gemini

openrouter.ai/deepseek/deepseek-v3.2-speciale

Campus Portal:Tata... | jobs | WORK | iktc quick | Drop me in Bookma... | ruby automation | TRADE | Automation | Java simple | Recipes | Gmail | YouTube | Maps | All Bookmarks

OpenRouter | Search | / | Fusion | Models | Chat | Rankings | Apps | Enterprise | Pricing | Docs | Sign Up

DeepSeek-V3.2-Speciale is a high-compute variant of DeepSeek-V3.2 optimized for maximum reasoning and agentic performance. It builds on DeepSeek Sparse Attention (DSA) for efficient long-context processing, then scales post-training reinforcement learning to push capability beyond the

Model weights

Overview | Playground | Providers | Performance | Pricing | Apps | Activity | Uptime | API

Providers for DeepSeek V3.2 Speciale

OpenRouter routes requests to the best providers that are able to handle your prompt size and parameters, with fallbacks to maximize uptime.

AlibabaCloud

Latency 2.88s | Throughput 26tps | Uptime

Total Context	Max Output	Input Price	Output Price	Cache Read
163.8K	163.8K	\$0.287 /K Tokens	\$0.431 /K Tokens	\$0.058 /K Tokens

Performance for DeepSeek V3.2 Speciale

Compare different providers across OpenRouter

All locations

Throughput | Latency | E2E Latency

openrouter.ai/deepseek/deepseek-v3.2-exp

DeepSeek-V3.2-Exp is an experimental large language model released by DeepSeek as an intermediate step between V3.1 and future architectures. It introduces DeepSeek Sparse Attention (DSA), a fine-grained sparse attention mechanism designed to improve training and inference efficiency in

Standard Model weights

Overview Playground Providers Performance Pricing Benchmarks Apps Activity Uptime API

Providers for DeepSeek V3.2 Exp

OpenRouter routes requests to the best providers that are able to handle your prompt size and parameters, with fallbacks to maximize uptime.

Sort by

Provider	Latency	Throughput	Uptime
NovitaAI	1.35s	15tps	100%
SiliconFlow	6.10s	7tps	100%

Model 6: Grok-provider-Xai-

Good in agentic work

openrouter.ai/provider/xai

Feb 8 Feb 18 Feb 28 Mar 10 Mar 20 Mar 30 Apr 9 Apr 19 Apr 29

xAI: Grok 4.3

Grok 4.3 is a reasoning model from xAI. It accepts text and image inputs with text output, and is suited for agentic workflows, instruction-following tasks, and applications requiring high factual accuracy. Reasoning...

by xAI · Apr 30, 2026 · 1M context · \$1.25/M input tokens · \$2.50/M output tokens

xAI: Grok 4.20 Multi-Agent

Grok 4.20 Multi-Agent is a variant of xAI's Grok 4.20 designed for collaborative, agent-based workflows. Multiple agents operate in parallel to conduct deep research, coordinate tool use, and synthesize informati...

by xAI · Mar 31, 2026 · 2M context · \$2/M input tokens · \$6/M output tokens

xAI: Grok 4.20

Grok 4.20 is a reasoning model from xAI with industry-leading speed and agentic tool calling capabilities. It combines the lowest hallucination rate on the market with strict prompt adherence, delivering consistently ...

by xAI · Mar 31, 2026 · 2M context · \$1.25/M input tokens · \$2.50/M output tokens

xAI: Grok 4.1 Fast


Going away May 15, 2026

Grok 4.1 Fast is xAI's best agentic tool calling model that shines in real-world use cases like customer support and deep research. Reasoning can be enabled/disabled using the 'reasoning' ...


by xAI · Nov 19, 2025 · 2M context · \$0.20/M input tokens · \$0.50/M output tokens

xAI: Grok 4 Fast

Model 7: Mistral-Provider-Mistral AI




Model 8:MiMo -provider-Xiaomi-Multimodel perception with agentic capability



Work : 2

Transformers concept – vector embedding and self attention mechanism

1.



2.


Transformers are language models

All the Transformer models mentioned above (GPT, BERT, T5, etc.) have been trained as *language models*. This means they have been trained on large amounts of raw text in a self-supervised fashion.

Self-supervised learning is a type of training in which the objective is automatically computed from the inputs of the model. That means that humans are not needed to label the data!

This type of model develops a statistical understanding of the language it has been trained on, but it's less useful for specific practical tasks. Because of this, the general pretrained model then goes through a process called *transfer learning* or *fine-tuning*. During this process, the model is fine-tuned in a supervised way — that is, using human-annotated labels — on a given task.

An example of a task is predicting the next word in a sentence having read the n previous words. This is called *causal language modeling* because the output depends on the past and present inputs, but not the future ones.




Another example is *masked language modeling*, in which the model predicts a masked word in the sentence.

3.

Text generation

Text generation involves creating coherent and contextually relevant text based on a prompt or input.

GPT-2 is a decoder-only model pretrained on a large amount of text. It can generate convincing (though not always true!) text given a prompt and complete other NLP tasks like question answering despite not being explicitly trained to.



1. GPT-2 uses byte pair encoding (BPE) to tokenize words and generate a token embedding. Positional encodings are added to the token embeddings to indicate the position of each token in the sequence. The input embeddings are passed through multiple decoder blocks to output some final hidden state. Within each decoder block, GPT-2 uses a *masked self-attention* layer which means GPT-2 can't attend to future tokens. It is only allowed to attend to tokens on the left. This is different from BERT's [mask] token because, in masked self-attention, an attention mask is used to set the score to 0 for future